

程序代写代做 CS编程辅导

SEC204



Computer architectures and low level programming

WeChat: cstutorcs

Assignment Project Exam Help

Dr. Vasilios Kelefouras

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

Email: [v.kelefouras@plymouth.ac.uk](mailto:v.kelefouras@plymouth.ac.uk)

Website: <https://www.plymouth.ac.uk/staff/vasilios-kelefouras>

<https://tutorcs.com>

# Outline

程序代写代做 CS编程辅导

2



- Different ways of writing assembly code
- Using intrinsic functions in C++
- Writing C/C++ programs using Intel SSE intrinsics
- Writing C/C++ programs using Intel AVX intrinsics

WeChat: cstutorcs

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

# Different ways of writing assembly

程序代写代做 CS编程辅导

3



1. Writing an entire program in assembly

2. Using inline assembly in C/C++

3. **Using intrinsic functions in C/C++**

- highly recommended, much easier and safer
- All the compilers support intrinsic functions
- An intrinsic function is equivalent to an assembly instruction
- **Mixes the good things of C++ (development time, portability, maintainability etc) with the good things of assembly (execution time)**

WeChat: tutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

□ C and C++ are the most easily combined languages with assembly code

# Different ways of writing assembly

## Using intrinsic functions in C/C++

程序代写代做 CS编程辅导

4



### □ Main advantages

- ▣ Classes, if conditionals and functions are very easy to implement
- ▣ Portability to almost all X86 architectures
- ▣ Compatibility with different compilers

WeChat: cstutorcs

### □ Main disadvantages

- ▣ Not all assembly instructions have intrinsic function equivalents
- ▣ Unskilled use of intrinsic functions can make the code less efficient than simple C++ code

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

# Using intrinsic functions in C/C++

程序代写代做 CS编程辅导

5



- For the rest of this lecture will be learning how to use intrinsic functions in C/C++

- Normally, “90% of a program’s execution time is spent in executing 10% of the code” - **loops**

WeChat: cstutorcs

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

- ▣ What programmers normally do to improve performance is to analyze the code and find the computationally intensive functions
  - Then optimize those instead of the whole program
  - This saves time and money

- ▣ **Rewriting loop kernels in C++ using SIMD intrinsics is an excellent choice**

<https://tutorcs.com>

- Compilers vectorize the code (not always) but manually using SIMD intrinsics can really boost performance

# Single Instruction Multiple Data (SIMD) –

程序代写代做CS编程辅导

6



```
for(i = 0; i <= MAX; i++)  
    c[i] = a[i] + b[i];
```

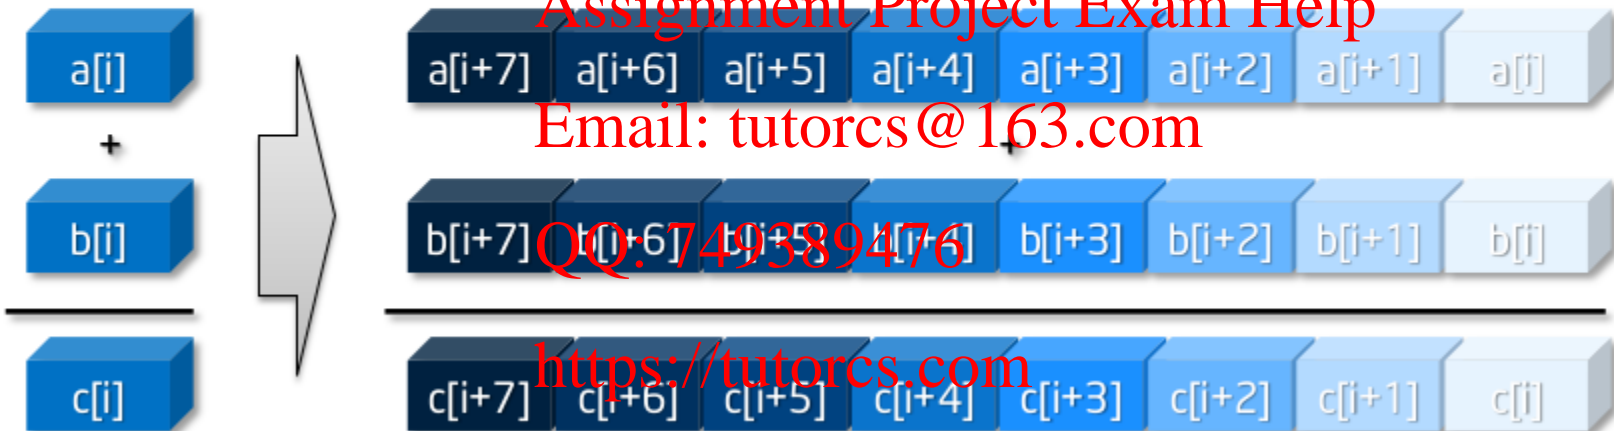
WeChat: cstutorcs

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>



# Vectorization on Arm Cortex series

NEON technology  
程序代写代做CS编程辅导

7

- Arm Neon technology is a parallel SIMD architecture extension for the Arm Cortex-A series and R52 processors



- ▣ 128-bit wide

- ▣ They are widely used in embedded systems

WeChat: cstutorcs

- Neon instructions allow up to:

- ▣ 16x8-bit, 8x16-bit, 4x32-bit, 2x64-bit integer operations

- ▣ 8x16-bit, 4x32-bit, 2x64-bit floating-point operations

QQ: 749389476

<https://tutorcs.com>

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

# Vectorization on Intel Processors

程序代写代做 CS编程辅导



WeChat: cstutorcs

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

## □ Intel MMX technology (old usage nowadays)

- 8 mmx registers of 64 bits
- extension of the floating point registers
- can be handled as 8 8-bit, 4 16-bit, 2 32-bit and 1 64-bit, operations
- An entire L1 cache line is loaded to the RF in 1-3 cycles

## □ Intel SSE technology

- 8/16 xmm registers of 128 bit (32-bit architectures support 8 registers only)
- Can be handled from 16 8-bit to 1 128-bit operations
- An entire L1 cache line is loaded to the RF in 1-3 cycles

## □ Intel AVX technology

- 8/16 ymm registers of 256 bit (32-bit architectures support 8 registers only)
- Can be handled from 32 8-bit to 1 256-bit operations

## □ Intel AVX-512 technology

- 32 ZMM 512-bit registers



# Vectorization on Intel Processors (2)

程序代写代做 CS编程辅导

9



SSE and AVX-128 types

WeChat: cstutorcs

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

AVX-256 types



# Vectorization on Intel Processors (3)

程序代写代做 CS编程辅导



- The developer can use either SSE or AVX or both
  - ▣ AVX instructions improve performance for data parallel input
  - ▣ SSE instructions are preferred for less data parallel algorithms
- Vector instructions work only for data that they are written in consecutive main memory addresses
- Aligned load/store instructions are faster than the no aligned ones.
- memory and arithmetical instructions are executed in parallel

WeChat: estutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

- **All the Intel intrinsics can be found here :**

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/#>

# Basic SSE Instructions (1)

程序代写代做 CS编程辅导

11



- `__m128 _mm_load_ps` - Loads four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_loadu_ps` - Loads four SP FP values. The address need not be 16-byte-aligned

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

WeChat: cstutorcs

Aligned load

Assignment Project Exam Help

Email: tutorcs@163.com

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

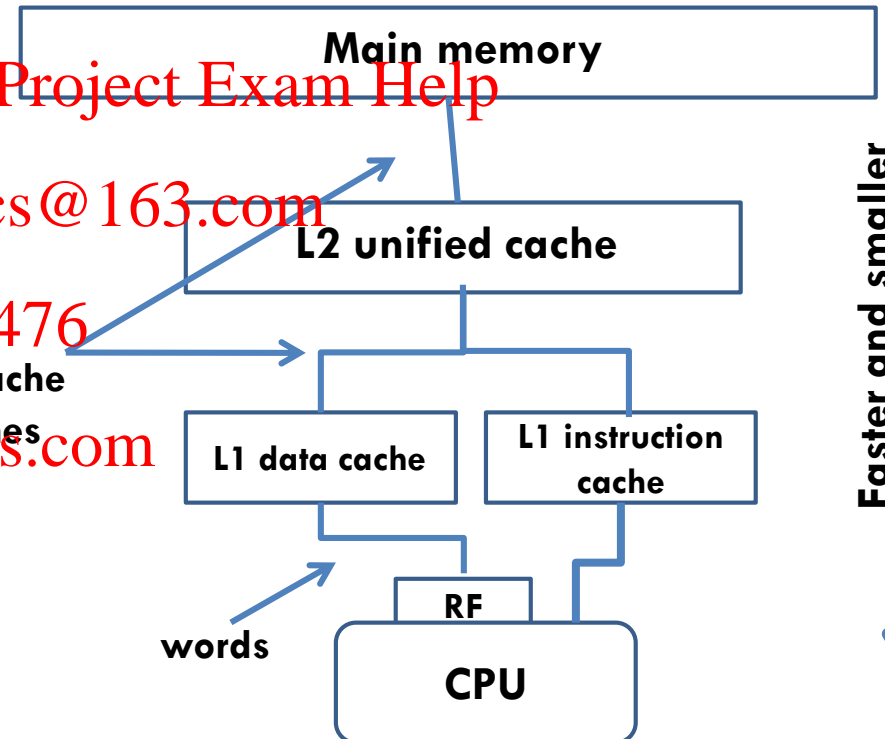
Misaligned load

https://tutorcs.com

L1

A[1]	A[2]	A[3]	A[4]
A[5]	A[6]	A[7]	A[8]
....			

Misaligned load



# Basic SSE Instructions (2)

## 程序代写代做 CS编程辅导

12



- `__m128 _mm_load_ps` - Loads four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_loadu_ps` - Loads four SP FP values. The address need not be 16-byte-aligned

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

`float A[N] __attribute__((aligned(16)));`  
WeChat: cstutorcs

Aligned load Main Memory  
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Misaligned load

Modulo (Address, 16) = 0

# Basic SSE Instructions (3)

程序代写代做 CS编程辅导

13



- `__m128 _mm_store_ps` – Stores four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_storeu_ps` – Stores four SP FP values. The address need not be 16-byte-aligned
- `__m128 _mm_mul_ps` – Multiplies the four SP FP values of a and b
- `__m128 _mm_mul_ss` – Multiplies the lower SP FP values of a and b; the upper 3 SP FP values are passed through from a.

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutores@163.com

`XMM1 = _mm_mul_ss(XMM1, XMM0)`      `XMM1 = _mm_mul_ps(XMM1, XMM0)`

	127	95	63	31	0
XMM0	4.0	3.0	2.0	1.0	
XMM1	5.0	5.0	5.0	5.0	
XMM1	4.0	3.0	2.0	5.0	

	127	95	63	31	0
XMM0	4.0	3.0	2.0	1.0	
XMM1	5.0	5.0	5.0	5.0	
XMM1	20.0	15.0	10.0	5.0	

https://tutorcs.com

QQ: 749389476

# Basic SSE Instructions (4)

## 程序代写代做 CS编程辅导

14



- `__m128 _mm_unpacklo_ps (__m128 a, __m128 b)` - Selects and interleaves the upper two SP FP values from a and b.
- `__m128 _mm_unpackhi_ps (__m128 a, __m128 b)` - Selects and interleaves the lower two SP FP values from a and b.

WeChat: cstutorcs

`XMM0=_mm_unpacklo_ps (XMM0, XMM1)`

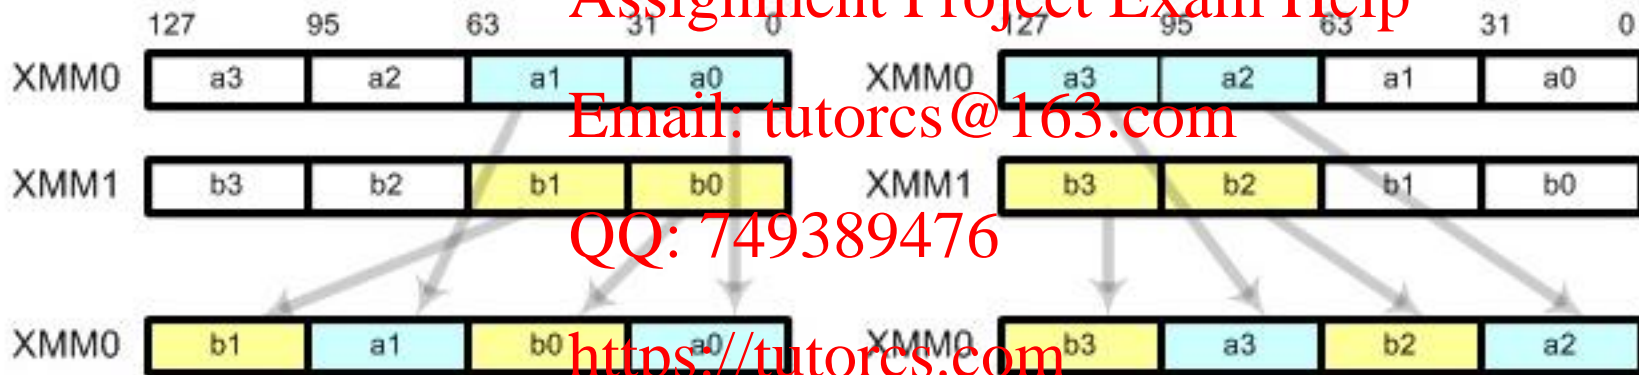
`XMM0=_mm_unpackhi_ps (XMM0, XMM1)`

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

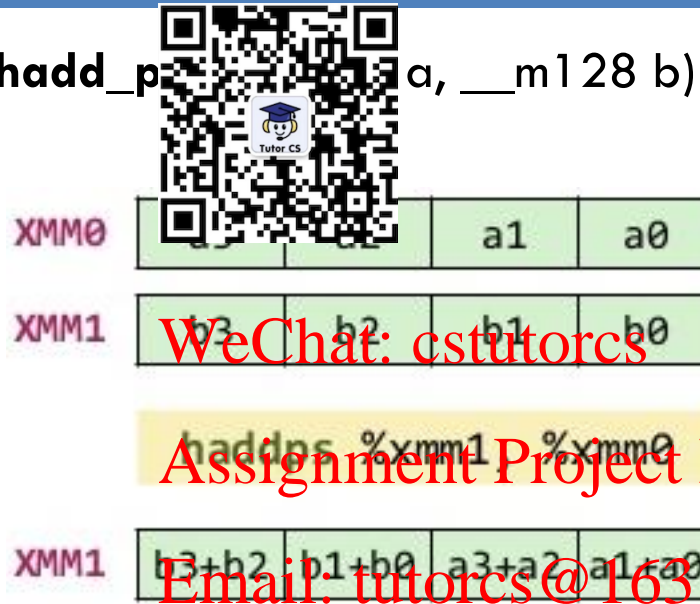


# Basic SSE Instructions (5)

程序代写代做 CS编程辅导

15

- `__m128 _mm_hadd_ps (__m128 a, __m128 b)` - Adds adjacent vector elements



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

- `void _mm_store_ss (float * p, __m128 a)` - Stores the lower SP FP value

<https://tutorcs.com>

# Case Study

## MVA using SSE technology

16



```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=0; i<N; i++){
```

WeChat: cstutorcs

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

```
for (j=0; j<N; j+=4){  
  num0=_mm_load_ps( &A[i][j] );  
  num1=_mm_load_ps(X + j );  
  num3=_mm_fmadd_ps(num0,num1,num3);  
  num4=_mm_hadd_ps(num3, num3);  
  num4=_mm_hadd_ps(num4, num4);  
  _mm_store_ss((float *)Y+i, num4);  
}
```



# MVM with SSE (2)

```
for (i=0; i!=N; i++){
```

```
num3=_mm_setzero_ps();
```



程序代写代做CS编程辅导




```
for (j=0; j!=N; j+=4){
```

```
num0=_mm_load_ps( &A[i][j] );
```

```
num1=_mm_load_ps(X + j );
```

```
num3=_mm_fmadd_ps(num0,
```



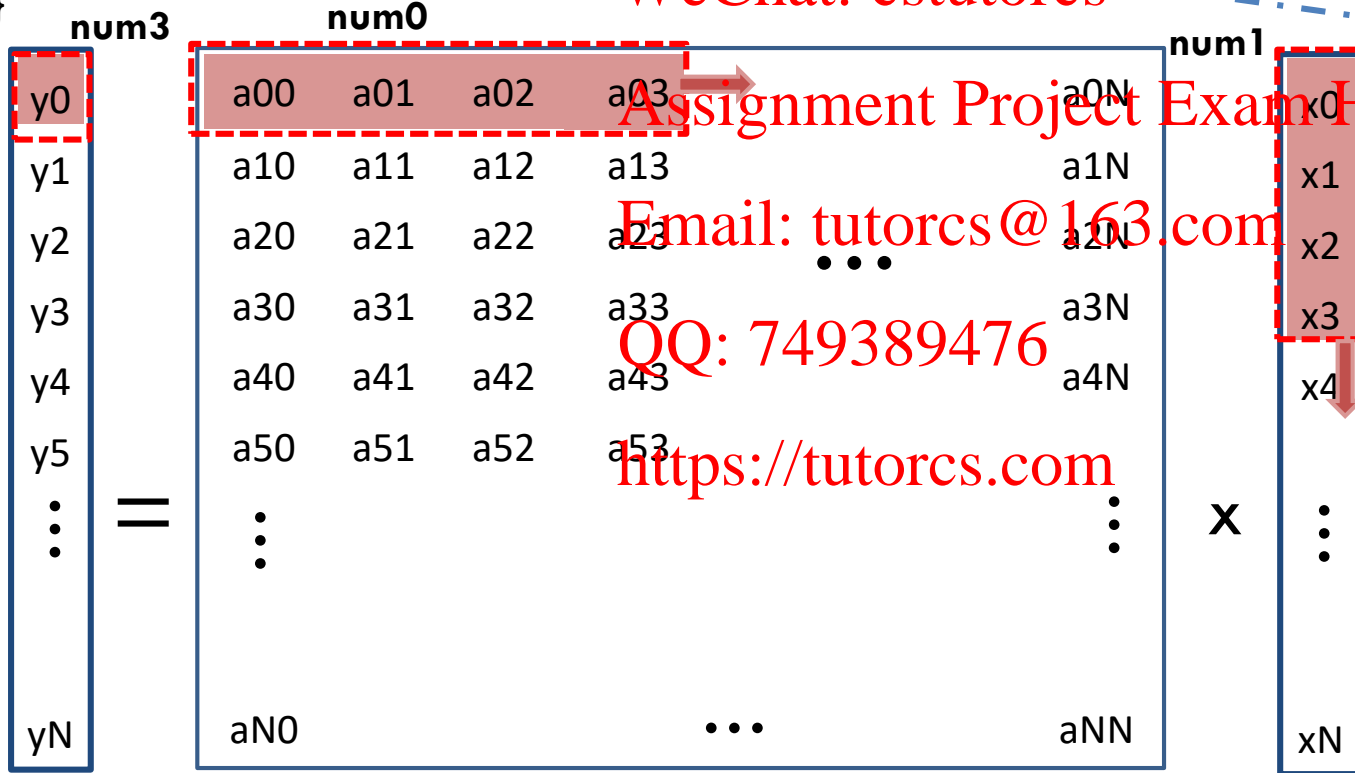
WeChat: cstutorcs

```
num3=_mm_hadd_ps(num3, num3);
```

```
num3=_mm_hadd_ps(num3, num3);
```

```
_mm_store_ss((float *)Y+i, num3);
```

```
}
```



This part of code adds the four values of num3 and stores the result into Y[i]

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

```

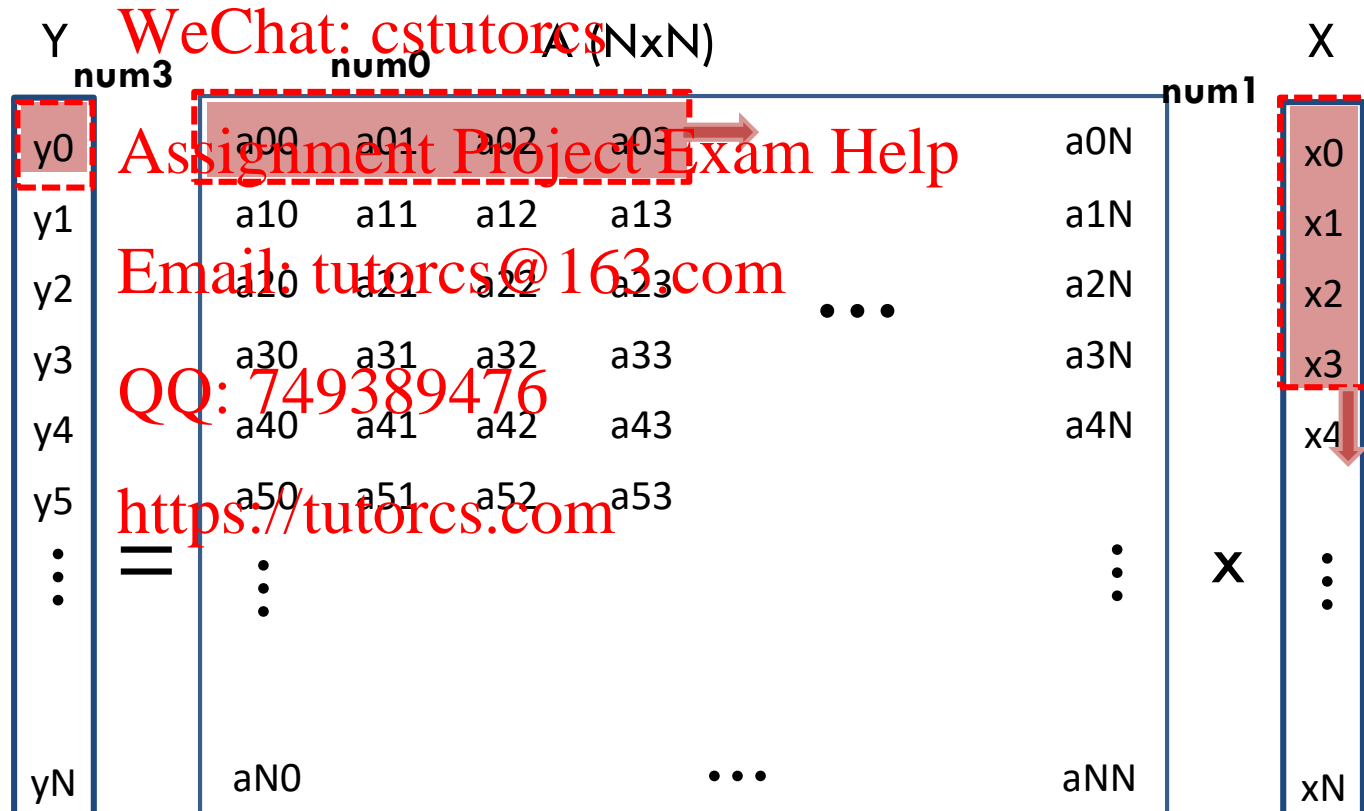
....
}
num3=_mm_hadd_ps(num3, num3);
num3=_mm_hadd_ps(num3, num3);
_mm_store_ss((float *)Y+i, num3);
}

```

程序代写代做 CS编程辅导



- After j loop finishes its execution, num3 contains the output data of Y[i]
- num3=[ya, yb, yc, yd]** where **ya=yb+yc+yd**
- after the 1<sup>st</sup> hadd -> num3=[yc+yd, ya+yb, yc+yd]
- after the 2<sup>nd</sup> hadd -> num3=[ya+yb+yc+yd, ya+yb+yc+yd, ya+yb+yc+yd, ya+yb+yc+yd]



# Case Study

## MVM using AVX technology

19



```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        Y[i] += A[i][j] * X[j];
```

```
for (i=0; i!=N; i++)  
    for (j=0; j!=N; j++) {  
        ymm0= _mm256_setzero_ps();  
        for (k=0; k!=N; k+=8) {  
            ymm1=_mm256_load_ps( A + N*i + k);  
            ymm2=_mm256_load_ps( Btrans + N*j + k);  
            ymm0=_mm256_fmadd_ps(ymm1, ymm2, ymm0);  
        }  
        ymm2 = _mm256_permute2f128_ps(ymm0, ymm0, 1);  
        ymm0 = _mm256_add_ps(ymm0, ymm2);  
        ymm0 = _mm256_hadd_ps(ymm0, ymm0);  
        ymm0 = _mm256_hadd_ps(ymm0, ymm0);  
        _mm_store_ss((float *) C + N*i + j,  
                     _mm256_extractf128_ps(ymm0, 0));  
    }  
}
```

WhatsApp: [tutorcs](https://tutorcs.com)

Assignment Project Exam Help

Email: [tutorcs@163.com](mailto:tutorcs@163.com)

QQ: 749389476

<https://tutorcs.com>

for (i=0; i < n; i++) {

if ( x[i] > 2 || x[i] < -2 )

a[i]+=x[i]; }

What about if-conditions on SSE ?

程序代写代做 CS编程辅导



```
const __m128 P2f = _mm_set_ps(2, 2, 2, 2);
const __m128 M2f = _mm_set_ps(-2, -2, -2, -2);
for (int i = 0; i < n; i++) {
    __m128 xv = _mm_load_ps(x + i);
    __m128 av = _mm_load_ps(a + i);

    __m128 c1v = _mm_cmpgt_ps(xv, P2f);
    __m128 c2v = _mm_cmplt_ps(xv, M2f);

    __m128 cv = _mm_or_ps(c1v, c2v);

    xv = _mm_and_ps(xv, cv);

    av = _mm_add_ps(av, xv);

    _mm_store_ps(a + i, av);
}
```



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

https://tutorcs.com

2	2	2	2
---	---	---	---

-2	-2	-2	-2
----	----	----	----

5	-3	0	1
---	----	---	---

a[i]	a[i+1]	a[i+2]	a[i+3]
------	--------	--------	--------

1	0	0	0
---	---	---	---

0	1	0	0
---	---	---	---

1	1	0	0
---	---	---	---

x[i]	x[i+1]	0	0
------	--------	---	---

a[i]	a[i+1]	a[i+2]	a[i+3]
+	+	+	+
x[i]	x[i+1]	0	0